

Course Overview

Python for AI-Driven Automation and Business Data Science

Course materials

Contents

- 1 The big picture
- 2 Structure
- 3 Learning paths
- 4 How to study
- 5 What you'll build
- 6 Getting started


Outline

- 1 The big picture
- 2 Structure
- 3 Learning paths
- 4 How to study
- 5 What you'll build
- 6 Getting started

The course in one diagram

Python for AI-Driven Automation and Business Data Science


38 notebooks · 10 modules · self-paced — spiral order



— runs entirely offline · every exercise has a worked solution · every notebook executes end-to-end —

Foundations of every page: read prose → run code → tweak → predict → re-run

Where this course sits



Three skills, taught together

1. Python fluency

Read, write, and modify clean Python without friction. Functions, modules, error handling, defensive parsing.

2. Working with data professionally

pandas, NumPy, matplotlib, scikit-learn, statistics, time-series forecasting. Real datasets, not toy ones.

3. AI-driven automation

Calling LLMs as functions. Structured output. Retrieval-augmented generation. Tools and agents. Document processing. *Plus* the engineering wrapper around all of it.

Outline

- 1 The big picture
- 2 Structure**
- 3 Learning paths
- 4 How to study
- 5 What you'll build
- 6 Getting started

Sixteen modules, forty-six lessons (+ 13 appendices)


#	Module	Notebooks
0	Onboarding	master onboarding + overview + see-it-work
1	Foundations	NB 1–6 — Python you can read without friction
2	Data science	NB 7–11 — pandas, NumPy, plots, stats, time series
3	Real-world I/O	NB 12–13 — HTTP, SQL, validation
4	Machine learning	NB 14–16 — sklearn, model evaluation, features
5	Industry applications	NB 17–20 — churn value, fraud, segmentation, demand
6	AI engineering	NB 21–26 — LLM patterns, RAG, agents, evaluation
7	Building AI POCs	NB 27–30 — setup, three POCs, RAG deep dive, agentic AI
8	Agents, tools & MCP	NB 31–34 — architectures, robust tools, MCP, multi-agent
9	NLP	NB 35–37 — topic modeling, sentiment
10	DeepTab (optional)	NB 38 — deep learning for tabular data
11	Production	NB 39–40 — packaging, scheduling, config & secrets
12	CI/CD & deployment	markdown module — Docker, GitHub Actions, HTTPS
13	Capstones	NB 41–42 — analytics + AI assistant
14	Business AI	NB 43–46 — strategy, architecture, BPM, governance
15	Django	markdown module — serve a model behind a web UI

Every module folder has its own README

Goals, prerequisites, what you'll build, the notebooks in order. Read it before you open the first notebook of the module.

Module dependencies — what comes before what

Module dependency graph — what depends on what




Edges flow downstream. Module N is safe to start once every upstream module is comfortable.

Outline

- 1 The big picture
- 2 Structure
- 3 Learning paths**
- 4 How to study
- 5 What you'll build
- 6 Getting started

Pick the path that fits you

Pick the learning path that fits you



Each dot is one module. Solid line = modules you'd touch on that path.

When to pick which path

Complete beginner

Never written code before. Take all modules in **spiral order** (see it work first, then the skills, then the builds). Time-budget: **about 115 hours** — the interactive estimator in 00b computes yours.

Analyst who knows Excel and SQL


Skim Module 1 (Python basics); Module 3's SQL will feel familiar. Focus on Modules 2 (data science), 4 (ML), 5 (industry applications), 13 (capstone). Time-budget: **about 43 hours**.

ML practitioner adding AI engineering

Skip foundations and data science. Focus on Modules 6 (AI engineering), 8 (agents & MCP), 11 (production), 13 (capstone). Time-budget: **about 34 hours**.

Eleven weeks, self-paced

A self-paced 11-week schedule



= 8-11 hours of focused study per week · more time → faster, less time → split a week in two


Got more time? Move faster. Got less? Split one week into two.

Outline

- 1 The big picture
- 2 Structure
- 3 Learning paths
- 4 How to study**
- 5 What you'll build
- 6 Getting started

The five-step study loop

The 5-step study loop — apply to every notebook



Apply this loop to *every* notebook. Genuine struggle beats passive reading.

Six visual markers you'll see throughout

Marker	Meaning
(Tip)	A small “stop and notice” callout.
(Intuition)	The mental model behind the syntax.
(Pitfall)	A bug or anti-pattern that bites if you're not warned.
(Exercises)	Hands-on tasks. Try <i>before</i> peeking at the solution.
(Bonus)	One larger applied task per notebook.
(Debug me)	A cell that <i>intentionally</i> contains a bug for you to find.
(Next step)	Always points to the next notebook in your path.

Once you internalise the markers they become road signs, not decoration.

Outline

- 1 The big picture
- 2 Structure
- 3 Learning paths
- 4 How to study
- 5 What you'll build**
- 6 Getting started

Six small artefacts you ship along the way

Six small artefacts you'll ship by the end

KPI snapshot

An AI-bot KPI report
in ~30 lines of Python

Live API ETL

Daily weather pulled
from Open-Meteo

SQL report

Channel summary written
in SQL + pandas

Forecast

3-month forecast
with Holt-Winters

Inbox triage

LLM that classifies +
summarises feedback

Scheduled job

Weekly report task
with retries + alerts

Plus two big capstones at the end

Capstone A — AI Support-Bot Analytics (NB 41)

Five channels $\times$ twelve months of operational data $\rightarrow$ 2 $\times$ 2 executive dashboard $\rightarrow$ regression demonstrating Simpson's paradox $\rightarrow$ five-bullet executive summary.

Capstone B — AI Customer-Feedback Assistant (NB 42)

Ingest free-form customer feedback $\rightarrow$ classify + extract via the MockLLM $\rightarrow$ ground answers in a knowledge base via RAG $\rightarrow$ orchestrate as a scheduled job with alerts. End-to-end AI feature.

Outline

- 1 The big picture
- 2 Structure
- 3 Learning paths
- 4 How to study
- 5 What you'll build
- 6 Getting started**

LLM providers — local and hosted, four options

Provider	Class	Best for
OpenAI	<code>OpenAILLM("gpt-4o-mini")</code>	Reliable default
Anthropic	<code>AnthropicLLM("claude-haiku")</code>	Long context, careful tone
Google	<code>GoogleLLM("gemini-2.0-flash")</code>	Cheapest at high volume
Ollama	<code>OllamaLLM("llama3.2:3b")</code>	Local – no internet, no key

Default = the offline MockLLM

Every AI notebook runs offline by default. Swapping in a real provider is a *one-line* change thanks to the unified interface in `llm_providers.py` at the repo root.

See the dedicated guide

Notebook `06_ai_engineering/A1_llm_providers_guide.ipynb` has the full setup walk-through, model recommendations, and cost estimates.

What you need to install

For Google Colab

Nothing. All required libraries are pre-installed. Open the notebook, hit Shift+Enter.

For local Jupyter

```
python -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
jupyter lab
```

Notebook 0 includes an environment-check cell you run to confirm everything works before you move on.

Now → open Notebook 0

One file to open. The rest unfolds from there.

`00_onboarding/00_master_onboarding.ipynb`

Three things that notebook does for you

- 1 Confirms your Python environment is correctly set up.
- 2 Explains the five-step study loop you'll use throughout.
- 3 Helps you pick the right learning path from the five options.

One last reminder before you start

Five-step study loop, in your pocket

Read the prose first → **Run** every code cell → **Try** the exercises before peeking → **Tweak** one number → **Predict** the new output before re-running.

When to use which marker

(Tip) information; (Intuition) mental model; (Pitfall) avoid this; (Exercise) hands-on practice; (Bonus) bigger applied task; (Debug me) intentional bug — find it.

When you're stuck for 15 minutes

Copy the full error into a search engine. 95% of beginner errors are well-documented. The course is designed to teach you to debug, not to spare you the experience.

Welcome.

38 notebooks await you.

Open Module 0 next.