

Digital Transformation & AI-Induced Change

Module 16 — Business AI in Practice · Notebook 49

Prof. Dr. Christoph Weisser

HSBI

Summer Semester 2026

2026-07-22

Digital Transformation & AI-Induced Change

Digital Transformation & AI-Induced Change
Module 16 — Business AI in Practice · Notebook 49

Prof. Dr. Christoph Weisser

HSBI

Summer Semester 2026

Contents

- 1 Why now
- 2 Tasks
- 3 Maturity
- 4 Strategies
- 5 Pitfalls
- 6 People

2026-07-22


Digital Transformation & AI-Induced Change

└ Contents

Contents
1 Why now
2 Tasks
3 Maturity
4 Strategies
5 Pitfalls
6 People

- 1Why now
- 2Tasks
- 3Maturity
- 4Strategies
- 5Pitfalls
- 6People

Why AI feels different *this* time



- **Cheap intelligence** — an LLM call costs cents and needs no in-house ML team.
- **Natural-language interface** — an organisational change, not just a technical one.


2026-07-22

Digital Transformation & AI-Induced Change

└ Why now

└ Why AI feels different *this* time

Why AI feels different *this* time



- **Cheap intelligence** — an LLM call costs cents and needs no in-house ML team.
- **Natural-language interface** — an organisational change, not just a technical one.

~8 min. Point: name what actually changed, so “AI is different now” stops being a slogan.

The two shifts that matter: **general** capability (one model, many tasks — no bespoke model per problem) and **near-zero marginal cost** (cents per call, no in-house ML team).

The under-appreciated line is the second bullet: a natural-language interface means the people who understand the work can now build with it directly. That is an organisational change, and it is why this lecture exists at all.

Ask the room: “Who here has automated something at work with an LLM without asking IT?” Usually several hands — that is stage 0/1 of the maturity model, arriving four slides early.

Stay honest: earlier waves also felt inevitable. The claim is that the economics changed, not that the hype is new.

- 1 Why now
- 2 Tasks
- 3 Maturity
- 4 Strategies
- 5 Pitfalls
- 6 People

The unit of change: tasks, not jobs

- A job is a *bundle of tasks*. AI absorbs *tasks*, rarely whole jobs.
- Enumerate the tasks; score each; start where the win is safe and easy.

Reframe the stakeholder question

“Will AI replace these 12 people?” → “Which of their *tasks* will AI absorb, and what do the people do with the freed-up time?”

2026-07-22

Digital Transformation & AI-Induced Change

└ Tasks

└ The unit of change: tasks, not jobs

~10 min — the most useful reframe in the lecture.

Point: jobs are bundles of tasks. AI absorbs *tasks*. Almost every overheated prediction comes from treating the job as the unit.

Do this live: pick a job the room knows — support agent, junior analyst, recruiter — and enumerate its tasks on the board. Eight to twelve is typical. Then mark which are patternable. Usually a third, and rarely the part people think of as the job.

That board exercise sets up the next slide’s 2×2 perfectly — you’re about to score exactly those tasks.

Handle the anxiety directly. The question “what do people do with the freed-up time?” is not rhetorical; if leadership has no answer, adoption fails regardless of model quality.


The unit of change: tasks, not jobs

- A job is a bundle of tasks. AI absorbs tasks, rarely whole jobs.
- Enumerate the tasks; score each; start where the win is safe and easy.

Reframe the stakeholder question

“Will AI replace these 12 people?” → “Which of their tasks will AI absorb, and what do the people do with the freed-up time?”

Where AI lands first — patternable × stakes




Begin in the **high-pattern** / **low-stakes** corner: cheap to build, cheap to be wrong.

2026-07-22

Digital Transformation & AI-Induced Change

└ Tasks

└ Where AI lands first — patternable × stakes



~**10 min.** Point: a placement tool they will reuse all course — it comes back in NB 52's readiness checklist (question 4).

The logic of the green corner: high-pattern means it'll probably work; low-stakes means a wrong answer is survivable. You want *both* for a first project, because you are also buying organisational trust, not just shipping a feature.

Put the tasks from the previous slide into the quadrants — with the room, not for them. Disagreements about where something sits are the valuable part; “high-stakes” is often assumed rather than argued.

The top-right is the trap: patternable *and* high-stakes (medical, credit, legal) is exactly where a confident demo is most seductive and most dangerous. Augment, don't automate — and note that the top-left never becomes an AI project at all.

Outline

- 1 Why now
- 2 Tasks
- 3 **Maturity**
- 4 Strategies
- 5 Pitfalls
- 6 People

2026-07-22


Digital Transformation & AI-Induced Change

└─ Maturity

└ Outline

- Why now
- Tasks
- Maturity**
- Strategies
- Pitfalls
- People

The five-stage AI maturity model



Most organisations sit between **1** and **2**. The model names the *next realistic step* — not the destination.


2026-07-22

Digital Transformation & AI-Induced Change

└ Maturity

└ The five-stage AI maturity model

The five-stage AI maturity model



Most organisations sit between 1 and 2. The model names the *next realistic step* — not the destination.

~**6 min** (this slide and the next together ~12).

Point: the model's use is *diagnostic* — it names the next realistic step. Nobody jumps from 1 to 4.

Stress the footnote: stage 4 is not the goal for most organisations. Getting from 1 to 2 — one feature genuinely live, with a named owner — is the hard, valuable transition, and it is mostly organisational rather than technical.

Ask the room: “Where is your employer, and what's the evidence?” Push for evidence; self-assessment runs about one stage optimistic.

Maturity — the signature behaviour of each stage

#	Stage	Signature behaviour
0	Unaware	Not in the strategy; staff quietly use ChatGPT on their phones.
1	Experimenting	A few pilots in a sandbox; results stay in slides.
2	Operationalising	One or two features live in production, owned by a real team.
3	Industrialising	A platform team; reusable infra; a feature ships in weeks.
4	AI-native	AI is in the <i>process design</i> , not bolted on.

2026-07-22

Digital Transformation & AI-Induced Change

└ Maturity

└ Maturity — the signature behaviour of each stage

Maturity — the signature behaviour of each stage

#	Stage	Signature behaviour
0	Unaware	Not in the strategy; staff quietly use ChatGPT on their phones.
1	Experimenting	A few pilots in a sandbox; results stay in slides.
2	Operationalising	One or two features live in production, owned by a real team.
3	Industrialising	A platform team; reusable infra; a feature ships in weeks.
4	AI-native	AI is in the process design, not bolted on.

~6 min. Point: the signature behaviours are how you diagnose a stage honestly — claims are unreliable, behaviours aren't.

Stage 0's line is the one that gets a laugh and then lands: “staff quietly use ChatGPT on their phones.” That is shadow AI — the work is already happening, ungoverned, and the strategy question is whether the organisation knows.

Stage 1 → 2 is the graveyard — “results stay in slides” is pilot purgatory, which returns as a named pitfall three slides later. Flag the connection now.

Read no more than two rows aloud; let them scan the rest.

Outline

- 1 Why now
- 2 Tasks
- 3 Maturity
- 4 Strategies**
- 5 Pitfalls
- 6 People


2026-07-22

Digital Transformation & AI-Induced Change

- └ Strategies
 - └ Outline

Outline	
1	Why now
2	Tasks
3	Maturity
4	Strategies
5	Pitfalls
6	People

Three change strategies



Intuition

Most real projects are **augmentation** in year one, drifting to **substitution** for the well-understood tasks as trust grows.

2026-07-22

Digital Transformation & AI-Induced Change Strategies

Three change strategies



~8 min. Point: three ambition levels, with change-management cost rising left to right.

The intuition block is the practical advice: **start with augmentation**, even where substitution would technically work. Augmentation gives you a human check while you learn the real error modes, and it is far easier to get accepted.

Reinvention is rare and worth being honest about: it means redesigning the process around what AI makes cheap, and it usually requires authority that a project team doesn't have. Most "reinvention" slides describe augmentation.

Connect to the 2×2: high-stakes quadrant → augmentation; low-stakes patternable → substitution is defensible.

Outline

- 1 Why now
- 2 Tasks
- 3 Maturity
- 4 Strategies
- 5 Pitfalls
- 6 People

2026-07-22

Digital Transformation & AI-Induced Change

└ Pitfalls

└ Outline

Outline

Why now

Tasks

Maturity

Strategies

Pitfalls

People

The five adoption pitfalls — and how to prevent them

Pitfall	Prevention
Solutionism	fix the process <i>before</i> you automate it
Pilot purgatory	define a production owner on day one
No monitoring	ship the eval dashboard <i>with</i> the feature
No change management	budget for adoption, training, incentives
Governance as afterthought	put ethics / safety in the first design review

“92 % accurate at launch” means nothing if nobody is watching the number six months later.

2026-07-22

Digital Transformation & AI-Induced Change

Pitfalls

The five adoption pitfalls — and how to prevent them

Pitfall	Prevention
Solutionism	fix the process <i>before</i> you automate it
Pilot purgatory	define a production owner on day one
No monitoring	ship the eval dashboard <i>with</i> the feature
No change management	budget for adoption, training, incentives
Governance as afterthought	put ethics / safety in the first design review

92 % accurate at launch means nothing if nobody is watching the number six months later.

~10 min. Point: every one of these is an *organisational* failure, not a modelling one. That’s the thread — none is fixed by a better model.

Solutionism deserves the most time: automating a broken process makes it broken faster. If nobody can describe the current process, that’s the project, not the AI.

The footnote is the line to land the section on — accuracy is a measurement at a moment, and models degrade as the world moves. This is exactly what Module 18 is about, and NB 32 builds the dashboard.

Ask the room: “Which of these five have you personally watched happen?” Pilot purgatory wins nearly every time.

Outline

- 1 Why now
- 2 Tasks
- 3 Maturity
- 4 Strategies
- 5 Pitfalls
- 6 People

2026-07-22

Digital Transformation & AI-Induced Change

- └ People
 - └ Outline

The human side

- Job descriptions, skill mix, and team shape begin changing the *day you ship*.
- Frame the change as augmentation, and give people a one-click override.

The cheapest trust intervention

Publish the AI’s **confidence score** next to every output. Users quickly learn when to trust it and when to double-check.

2026-07-22

Digital Transformation & AI-Induced Change

└ People

└ The human side

~8 min. Point: adoption is a people problem, and it starts the day you ship — not later.

The confidence-score tip is cheap and genuinely effective: it lets users build a calibrated mental model of when the system is reliable, instead of swinging between blind trust and blanket rejection. Pair it with the **one-click override** — people accept a system far more readily when they can overrule it, even if they rarely do.

This connects forward to NB 52’s invoice case, where trust grew from *observation* over weeks. You cannot argue people into trusting a model.

If the cohort is technical, this is the slide they will underrate. Say so.

The human side

- Job descriptions, skill mix, and team shape begin changing the *day you ship*.
- Frame the change as augmentation, and give people a one-click override.

The cheapest trust intervention

Publish the AI’s **confidence score** next to every output. Users quickly learn when to trust it and when to double-check.

Companion notebook

This is the conceptual half of the lecture

The hands-on judgement — task inventories, placing an organisation on the maturity model, the pre-mortem, the strategy-map 2×2 — lives in **Notebook 49**, with worked solutions.

- Prerequisite-free; the course’s *spiral* recommends it as an early read (the “why”).
- Next: **Notebook 50 — Architecture Patterns**.

2026-07-22

Digital Transformation & AI-Induced Change

└ People

└ Companion notebook

Companion notebook

This is the conceptual half of the lecture

The hands-on judgement — task inventories, placing an organisation on the maturity model, the pre-mortem, the strategy-map 2×2 — lives in **Notebook 49**, with worked solutions.

- Prerequisite-free; the course’s *spiral* recommends it as an early read (the “why”).
- Next: **Notebook 50 — Architecture Patterns**.

~4 min — close. Point: hand off to the notebook, which is where the judgement is actually practised.

Name the four exercises so they know what’s coming: task inventory, placing an organisation on the maturity model, the pre-mortem, and the strategy-map 2×2 — all with worked solutions.

The pre-mortem is worth selling: “it’s 12 months from now and the project failed — write the reasons.” It surfaces the pitfalls from two slides ago in the room’s own words.

Remind them this notebook is prerequisite-free — they can send it to a non-technical colleague, which is unusual for this course.